

Importing Libraries

```
In [2]: import numpy as np # linear algebra
import pandas as pd # for data processing like in this project csv loading
import matplotlib.pyplot as plt # for data visualization
import seaborn as sns # for data visualization i.e pairplot
from sklearn import metrics # for calculating rootmean square
```

Data Loading and Insites

```
In [3]: # reading the data

df = pd.read_csv('insurance.csv')

# checking the shape
print(df.shape)

(1338, 8)
```

```
In [4]: # checking data points
print(df.size)

10704
```

```
In [5]: # previewing 1st 5 dataset checking wether its loaded or not sucessfully
df.head()
```

```
Out[5]:
```

	age	sex	bmi	children	smoker	region	charges	insuranceclaim
0	19	0	27.900	0	1	3	16884.92400	1
1	18	1	33.770	1	0	2	1725.55230	1
2	28	1	33.000	3	0	2	4449.46200	0
3	33	1	22.705	0	0	1	21984.47061	0
4	32	1	28.880	0	0	1	3866.85520	1

```
In [ ]:
```

In [6]: *#description about data set*

```
df.describe()
```

Out[6]:

	age	sex	bmi	children	smoker	region	char
count	1338.000000	1338.000000	1338.000000	1338.000000	1338.000000	1338.000000	1338.000
mean	39.207025	0.505232	30.663397	1.094918	0.204783	1.515695	13270.422
std	14.049960	0.500160	6.098187	1.205493	0.403694	1.104885	12110.011
min	18.000000	0.000000	15.960000	0.000000	0.000000	0.000000	1121.873
25%	27.000000	0.000000	26.296250	0.000000	0.000000	1.000000	4740.287
50%	39.000000	1.000000	30.400000	1.000000	0.000000	2.000000	9382.033
75%	51.000000	1.000000	34.693750	2.000000	0.000000	2.000000	16639.912
max	64.000000	1.000000	53.130000	5.000000	1.000000	3.000000	63770.428

In [7]: *#cheking information about data*

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   age                   1338 non-null  int64  
1   sex                   1338 non-null  int64  
2   bmi                   1338 non-null  float64
3   children              1338 non-null  int64  
4   smoker               1338 non-null  int64  
5   region               1338 non-null  int64  
6   charges              1338 non-null  float64
7   insuranceclaim       1338 non-null  int64  
dtypes: float64(2), int64(6)
memory usage: 83.8 KB
```

In [8]: *# checking number of null value in this data*

```
df.isnull().sum()
```

Out[8]:

age	0
sex	0
bmi	0
children	0
smoker	0
region	0
charges	0
insuranceclaim	0

dtype: int64

```
In [9]: # checking if any null value is present or not
df.isnull().any()
```


```
Out[9]: age                False
sex                False
bmi               False
children          False
smoker            False
region            False
charges           False
insuranceclaim    False
dtype: bool
```

```
In [10]: # from this data we can get insites that :
# 1. data belongs to middle age people (mostly)
# 2. maximum age of any person is 64 where as minimum age is 18 only
# 3. maximum bmi is 53.13 which is a deep sign of obesity
# 4. there is no null value in this data
# 5. There are 676 male and 662 female
```

```
In [11]: # checking value count of male and female in data
df['sex'].value_counts()
```


```
Out[11]: 1    676
0    662
Name: sex, dtype: int64
```

```
In [12]: # plotting a bar graph showing about number of male and female  
df.sex.value_counts(normalize=False).plot.barh()  
plt.show()
```



```
In [13]: #pie chart: with Label and explode
mylables=["Male","Female"] # here label is "Male - is 1 where as Female - is 0
colors = ['orange', 'pink']
myexplode=[0.10,0]
size = [676, 662]
plt.pie(size,colors = colors,labels =mylables,explode = myexplode, shadow = True)
plt.title('PIE chart representing share of men and women in insurance data ')
plt.legend()
plt.show()
```

PIE chart representing share of men and women in insurance data




```
In [14]: # checking customer belonging
df['region'].value_counts()
```

```
Out[14]: 2    364
          3    325
          1    325
          0    324
          Name: region, dtype: int64
```

```
In [15]: #ploting a Countplot showing region
sns.countplot("region",data = df)
plt.show()
```

D:\Anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.


```
warnings.warn(
```



```
In [16]: #plotting a Countplot showing age
plt.figure(figsize = (10,5))
sns.countplot("age",data = df)
plt.show()
```


D:\Anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

warnings.warn(



In [17]: *# plotting a bar graph showing about region wise with labels grid and minor grid*

```
x = ['North-East(0)', 'North-West(1)', 'South-East(2)', 'South-West(3)']
size = [324, 325, 364, 325]
plt.figure(figsize = (15,7))
x_pos = [i for i, _ in enumerate(x)]
plt.bar(x_pos, size, color=['cyan', 'yellow', 'green', 'orange'])
plt.xlabel("Region")
plt.ylabel("Size")
plt.title("Number of Insurance Holder from All Over the region")
plt.xticks(x_pos, x)
# Turn on the grid
plt.minorticks_on()
plt.grid(which='major', linestyle='-', linewidth='0.5', color='black')
# Customize the minor grid
plt.grid(which='minor', linestyle=':', linewidth='0.5', color='grey')
plt.show()
```



In [18]: *# 6.people belonging residential area from northeast(0) are 324 person ; north*

```
In [19]: # checking children count
df['children'].value_counts()
```


```
Out[19]: 0    574
         1    324
         2    240
         3    157
         4     25
         5     18
         Name: children, dtype: int64
```



```
In [20]: #ploting a Countplot showing number of children
plt.figure(figsize = (10,5))
sns.countplot("children",data = df)
plt.show()
```


D:\Anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

warnings.warn(



```
In [21]: #pie chart: with Label and explode
plt.figure(figsize = (10,5))
mylables=["No Children","1 Child","2 Children","3 Children","4 Children","5 Ch
colors = ['orange','pink','purple','cyan','yellow','blue']
myexplode=[0.10,0,0,0,0,0]
size = [574, 324,240,157,25,18]
plt.pie(size,colors = colors,labels =mylables,explode = myexplode, shadow = Tr
plt.title('PIE chart representing share of person per children as per given dat
plt.legend()
plt.show()
```

PIE chart representing share of person per children as per given data




```
In [22]: # 7. count of people having no children are 574 ; with 1 children are 324 ; wi
```

```
In [23]: # checking number of smokers
df['smoker'].value_counts()
```

```
Out[23]: 0    1064
         1     274
         Name: smoker, dtype: int64
```

```
In [24]: #ploting a bar grap showing number of smoker
plt.figure(figsize = (10,5))
mylables=['Non-Smoker','Smoker']
colors = ['pink','Red']
myexplode=[0.10,0.10]
size = [1064,274]
plt.pie(size,colors = colors,labels =mylables,explode = myexplode, shadow = Tr

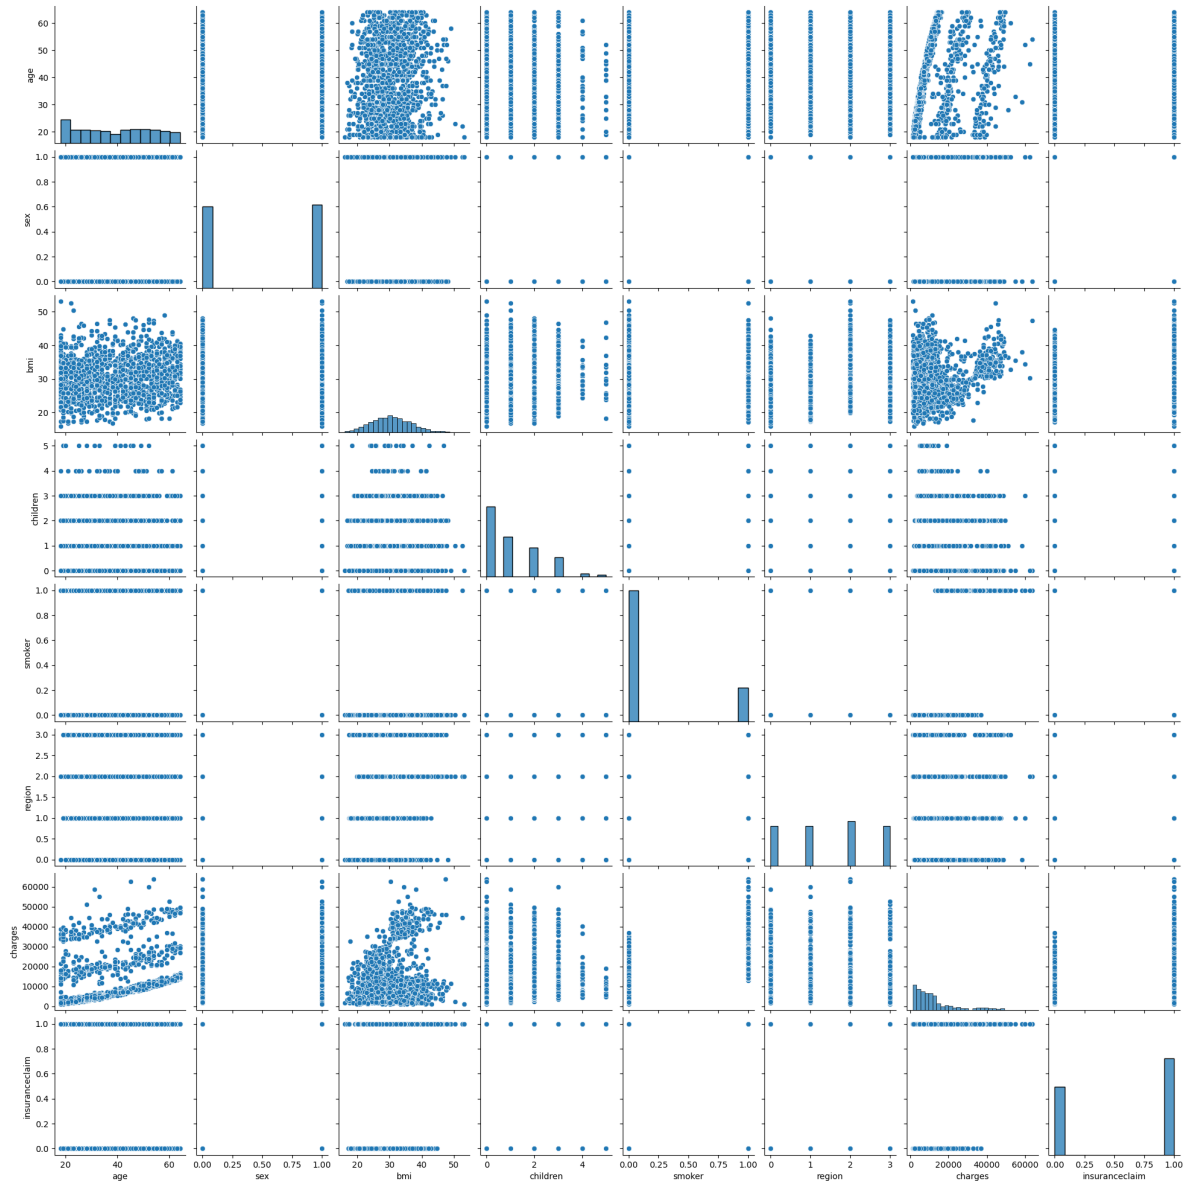
plt.show()
```



```
In [24]: # 8. count of non smokers are represented as 0 which is 1064 where as 274 peop
```

```
In [25]: # pairplot
sns.pairplot(df)
```

```
Out[25]: <seaborn.axisgrid.PairGrid at 0x1c696bf81c0>
```



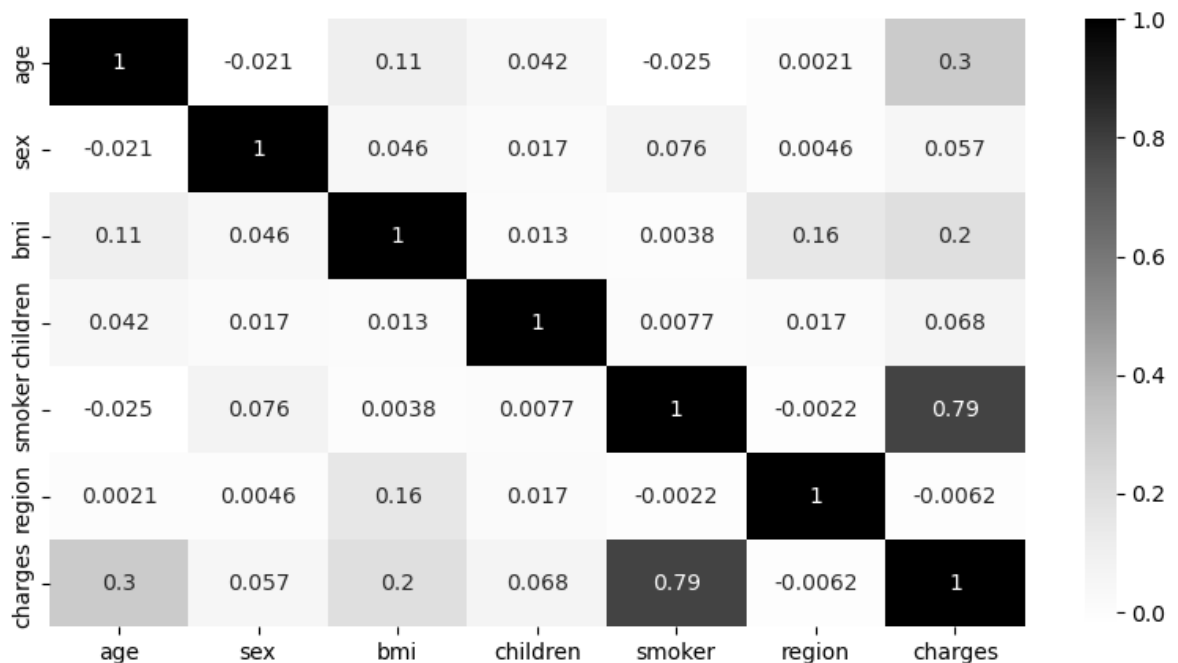
```
In [26]: # Correlation Between Different Feature
df[['age', 'sex', 'bmi', 'children', 'smoker', 'region', 'charges']].corr()
```

```
Out[26]:
```

	age	sex	bmi	children	smoker	region	charges
age	1.000000	-0.020856	0.109272	0.042469	-0.025019	0.002127	0.299008
sex	-0.020856	1.000000	0.046371	0.017163	0.076185	0.004588	0.057292
bmi	0.109272	0.046371	1.000000	0.012759	0.003750	0.157566	0.198341
children	0.042469	0.017163	0.012759	1.000000	0.007673	0.016569	0.067998
smoker	-0.025019	0.076185	0.003750	0.007673	1.000000	-0.002181	0.787251
region	0.002127	0.004588	0.157566	0.016569	-0.002181	1.000000	-0.006208
charges	0.299008	0.057292	0.198341	0.067998	0.787251	-0.006208	1.000000

```
In [27]: # Insite from this Correlation are : 1. Charges are Dependent Upon Age (Higher
```

```
In [28]: #plot the correlation matrix of salary, balance and age in data dataframe.
plt.figure(figsize = (10,5))
sns.heatmap(df[['age', 'sex', 'bmi', 'children', 'smoker', 'region', 'charges']
plt.show()
```



```
In [29]: df.columns
```

```
Out[29]: Index(['age', 'sex', 'bmi', 'children', 'smoker', 'region', 'charges',
               'insuranceclaim'],
              dtype='object')
```

In [30]: *# Insites From this Heat Map :- 1. Smoker Tends to Pay More Insurance Charges;*




In []:

```
In [31]: # Age vs Charges
# the more the age the more will be insurance charge (roughly estimated)

plt.figure(figsize = (10, 5))
sns.lineplot(x = 'age', y = 'charges', data = df)

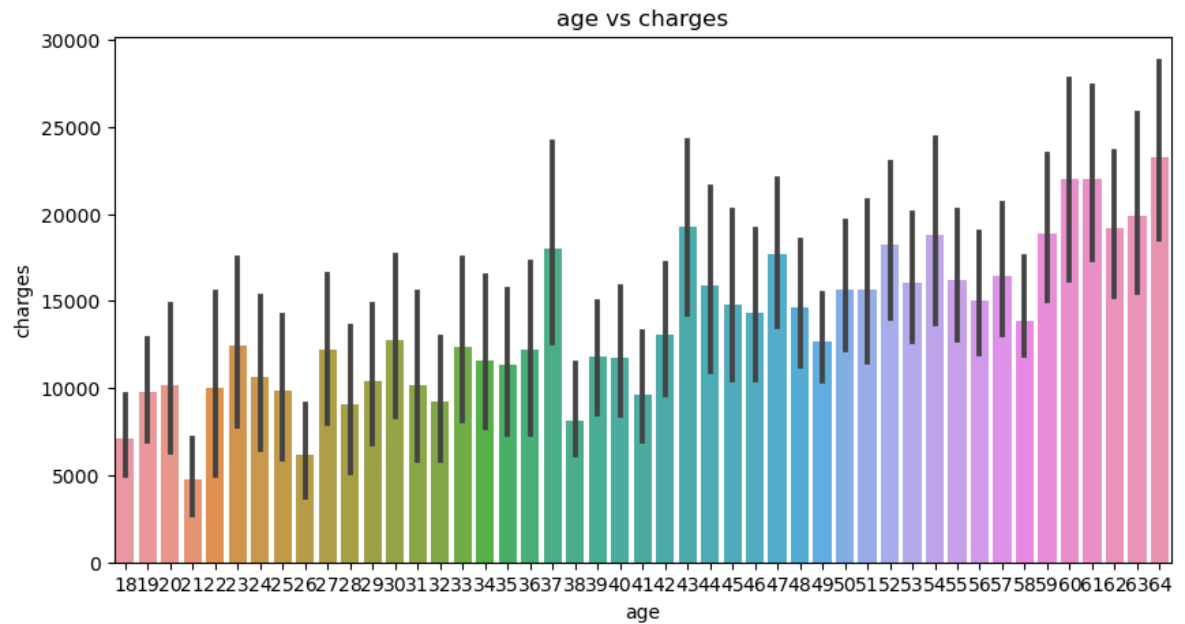
plt.title("Age vs Charges")
plt.show()
```



In [32]: *#box plot for age vs charge*

```
plt.figure(figsize = (10, 5))
sns.barplot(x = 'age', y = 'charges', data = df)

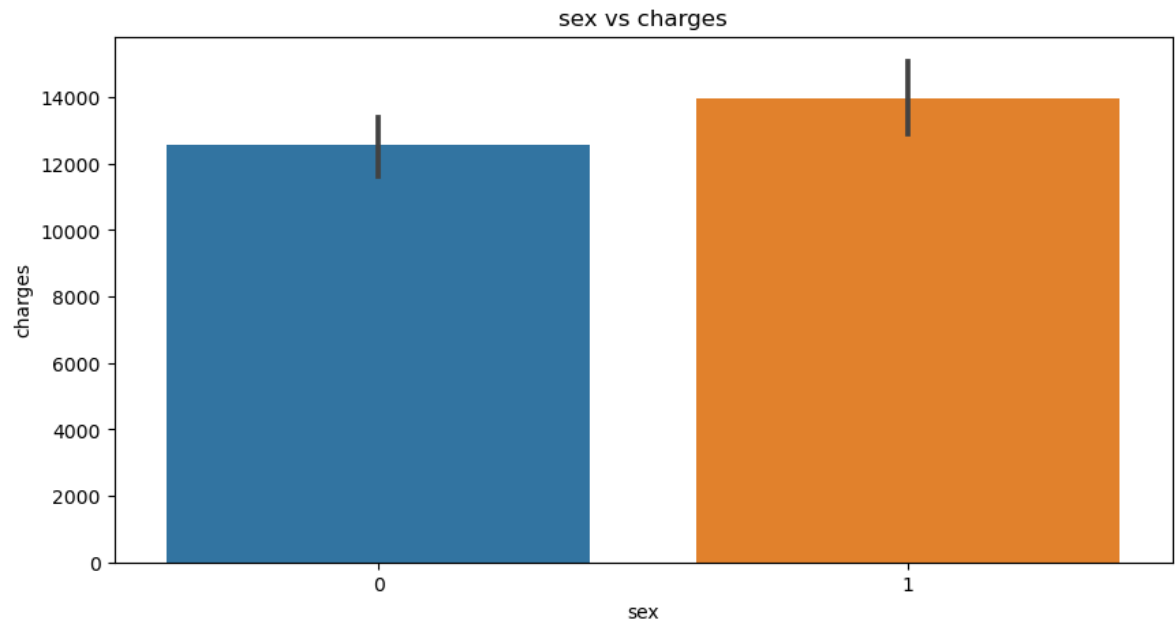
plt.title('age vs charges')
plt.show()
```



```
In [33]: #plot the box plot of sex and charges
# as 1 belongs to men : it shows that men are paying more insurance charges than women
#bar plot

plt.figure(figsize = (10, 5))
sns.barplot(x = 'sex', y = 'charges', data = df)


plt.title('sex vs charges')
plt.show()
```




```
In [34]: # children vs charges
# no. of children of a person has a weird dependency on insurance charge. i.e

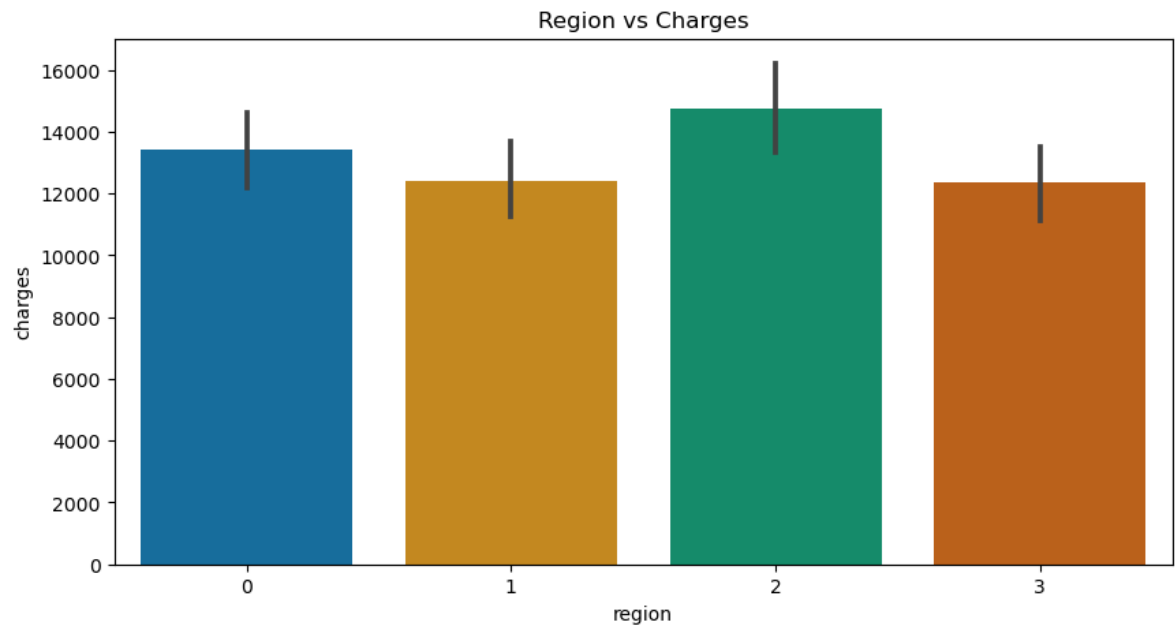
plt.figure(figsize = (10, 5))
sns.barplot(x = 'children', y = 'charges', data = df)

plt.title('CHILDREN VS CHARGES')
plt.show()
```



In [35]: *# region vs charges BAR GRAPh*


```
plt.figure(figsize = (10, 5))  
sns.barplot(x = 'region', y = 'charges', data = df, palette = 'colorblind')  
  
plt.title('Region vs Charges')  
plt.show()
```




In [36]: *# from the graph we can clearly state that region dont play any role in charge*

```
In [37]: # smoker vs charges
plt.figure(figsize = (10, 5))
sns.barplot(x = 'smoker', y = 'charges', data = df)

plt.title('SMOKERS VS CHARGES')
plt.show()
```




In [38]: *# from the graph where 0 represents non smoker and 1 represent smoker it is clear*



```
In [39]: # BMI vs charges
plt.figure(figsize = (20,8))
sns.barplot(x = 'bmi', y = 'charges', data = df)

plt.title('BMI VS CHARGES')
plt.show()
```



In [40]: *# From Graph we can conclude that higher the BMI more will be Insurance Premium*

Data Cleaning

In [41]: *# removing un required columns from the insurance data*
As from the above grph we can clearly state that region dont play any role in

```
df = df.drop('region', axis = 1)
```

In [42]: df.shape

Out[42]: (1338, 7)

In [43]: *#as earlier there was 10704 data point the new one has 9366 data point after removing region*
df.size

Out[43]: 9366

In [44]: *# seperate out features and target value from dataset*

```
X=df.drop(["insuranceclaim"],axis=1).values  
y=df["insuranceclaim"].values
```

In [45]: X.shape

Out[45]: (1338, 6)

In [46]: y.shape

Out[46]: (1338,)


Finding an Outlier

```
In [47]: #bmi outlier
```

```
sns.boxplot(df["bmi"])  
plt.show()
```

D:\Anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```



```
In [48]: # Finding Position of Outlier  
#position plot of outlier
```

```
print(np.where(df["bmi"]>45))
```

```
(array([ 116,  286,  292,  401,  438,  454,  543,  547,  549,  582,  660,  
        847,  860,  930,  941, 1024, 1047, 1088, 1131, 1317], dtype=int64),)
```


```
In [49]: # bmi can be more or less as per medical condition of person so no need to tre
```

In [50]: *#Charges outlier*

```
sns.boxplot(df["charges"])  
plt.show()
```

D:\Anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

warnings.warn(



In [51]: *# Charges can be More or Less as per required by insurance company*

Splitting Data (Training and Testing Data) and Importing Sklearn Modules

In [52]: *#splitting data into training and testing data set*

```
from sklearn.model_selection import train_test_split
```

```
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size = 0.20, random_
```

```
In [53]: print("X_train shape : " , X_train.shape)
print("X_test shape : ", X_test.shape)
print("y_train shape : " , y_train.shape)
print("y_test shape : ", y_test.shape)
```

```
X_train shape : (1070, 6)
X_test shape : (268, 6)
y_train shape : (1070,)
y_test shape : (268,)
```

Importing and Using Decision Tree (Supervised Learning) Algorithm

```
In [54]: from sklearn.tree import DecisionTreeClassifier
```

```
In [55]: # model

dtc = DecisionTreeClassifier()

#fitting

dtc.fit(X_train,y_train)
```

```
Out[55]: DecisionTreeClassifier()
```

```
In [56]: #predicting via Decision Tree Algorithm

y_pred=dtc.predict(X_test)

y_pred
```

```
Out[56]: array([0, 1, 1, 1, 1, 1, 0, 1, 1, 1, 0, 0, 0, 1, 1, 1, 1, 1, 0, 1, 0, 0,
 1, 1, 1, 1, 1, 0, 1, 1, 0, 1, 1, 0, 0, 1, 0, 1, 1, 1, 0, 1, 1, 0,
 0, 0, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 0, 1, 1, 1, 1, 0, 0, 0, 1, 0,
 0, 1, 0, 1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 1, 1, 0, 1, 1, 1, 0, 1,
 0, 0, 1, 1, 1, 0, 1, 0, 1, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0, 1,
 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 1, 1, 0, 0, 1, 1, 1, 1, 0,
 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 0, 1, 0, 0, 1, 1, 1, 1,
 1, 1, 1, 0, 1, 1, 0, 0, 1, 0, 0, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1, 1,
 1, 1, 0, 0, 1, 1, 0, 1, 0, 1, 0, 1, 1, 1, 0, 1, 0, 1, 1, 1, 0, 1,
 1, 0, 0, 1, 0, 0, 1, 1, 1, 0, 1, 0, 0, 0, 1, 1, 0, 1, 0, 1, 0, 1,
 1, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 1,
 1, 0, 0, 1, 0, 1, 0, 1, 1, 0, 0, 0, 1, 1, 1, 1, 0, 0, 0, 0, 1, 1,
 1, 0, 1, 1], dtype=int64)
```

In [57]: *#Calculating RMSE Root MEan Square Error*

```
rmse= np.sqrt(metrics.mean_squared_error(y_test,y_pred))  
print("Root Mean Square Error = ",rmse)
```

Root Mean Square Error = 0.16161498378886355

Checking Out Training and Testing Data Accuracy (Actual vs Predicted)

In [58]: *# compute accuracy on training set*

```
dtc_train= dtc.score(X_train,y_train)  
  
print("Training Data Accuracy by Decision Tree Algorithm is : " , dtc_train)  
  
# compute accuracy on testing set  
  
dtc_test= dtc.score(X_test,y_test)  
  
print("Testing Data Accuracy by Decision Tree Algorithm is : " , dtc_test)
```

Training Data Accuracy by Decision Tree Algorithm is : 1.0

Testing Data Accuracy by Decision Tree Algorithm is : 0.9738805970149254

In [59]: *# Here Training Accuracy is Greater than Testing Accuracy means our Model is Overfitting*

In [60]: *# calculating the mean squared error*
mse = np.mean((y_test - y_pred)**2, axis = None)
print("MSE :", mse)

Calculating the root mean squared error
rmse = np.sqrt(mse)
print("RMSE :", rmse)

MSE : 0.026119402985074626

RMSE : 0.16161498378886355

Using Hyperparameter Tuning for Decision Tree


```
In [61]: parameters = {"splitter" : ["best", "random"],
                        "max_depth" : [1,3,5,7,9,11],
                        "min_samples_leaf" : [1,2,3,4,5,6,7,8,9,10],
                        "min_weight_fraction_leaf": [0.2,0.3],
                        "max_features":["auto","log2","sqrt",None],
                        "max_leaf_nodes":[None,10,20,30]
                    }
```

```
In [62]: #using grid search cv

from sklearn.model_selection import GridSearchCV
```

```
In [63]: tuning_model = GridSearchCV(dtc,param_grid = parameters,
                                     scoring= "neg_mean_squared_error",
                                     cv=3,verbose=3)
```

```
In [ ]: tuning_model.fit(X,y)
```

```
In [82]: dtc.get_params().keys()
```

```
Out[82]: dict_keys(['ccp_alpha', 'class_weight', 'criterion', 'max_depth', 'max_features', 'max_leaf_nodes', 'min_impurity_decrease', 'min_samples_leaf', 'min_samples_split', 'min_weight_fraction_leaf', 'random_state', 'splitter'])
```

```
In [83]: #best parameters

tuning_model.best_params_
```

```
Out[83]: {'max_depth': 3,
          'max_features': None,
          'max_leaf_nodes': None,
          'min_samples_leaf': 1,
          'min_weight_fraction_leaf': 0.2,
          'splitter': 'best'}
```

```
In [84]: #using this type of hyper parameters to train our model once again

tuned_model=DecisionTreeClassifier(max_depth=3,min_samples_leaf=1,min_weight_fraction_leaf=0.2,splitter="best")
```

```
In [85]: #fitting model

tuned_model.fit(X_train,y_train)
```

```
Out[85]: DecisionTreeClassifier(max_depth=3, min_weight_fraction_leaf=0.2)
```

```
In [86]: #prediction
tuned_pred=tuned_model.predict(X_test)
```

```
In [87]: # compute accuracy on training set
tuned_model_train= tuned_model.score(X_train,y_train)
print("Training Data Accuracy by Decision Tree Tuned Algorithm is : " , tuned_
# compute accuracy on testing set
tuned_model_test= tuned_model.score(X_test,y_test)
print("Testing Data Accuracy by Decision Tree Algorithm is : " , tuned_model_t
```

Training Data Accuracy by Decision Tree Tuned Algorithm is : 0.8046728971962617
Testing Data Accuracy by Decision Tree Algorithm is : 0.8208955223880597

```
In [88]: #Calculating RMSE
rmse= np.sqrt(metrics.mean_squared_error(y_test,tuned_pred))
print("Root Mean Square Error = ",rmse)
```

Root Mean Square Error = 0.42320736951515897

```
In [89]: # calculating the mean squared error
mse = np.mean((y_test - y_pred)**2, axis = None)
print("MSE :", mse)
```

MSE : 0.1791044776119403

Importing and Using Logistic Regression

```
In [90]: from sklearn.metrics import accuracy_score, confusion_matrix # importing for er
from sklearn.linear_model import LogisticRegression # importing Logistic Regres
```

```
In [91]: # Logistic Regression model
logreg = LogisticRegression()
logreg.fit(X_train, y_train)
```

D:\Anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)
n_iter_i = _check_optimize_result(

```
Out[91]: LogisticRegression()
```

```
In [92]: y_pred = logreg.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
print("Accuracy: ", accuracy)
print("Confusion matrix: \n", conf_matrix)
```

Accuracy: 0.8208955223880597
Confusion matrix:
[[81 26]
 [22 139]]

```
In [93]: # compute accuracy on training set

logreg_train= logreg.score(X_train,y_train)

print("Training Data Accuracy by Logistics Regression Algorithm is : " ,logreg

# compute accuracy on testing set

logreg_test= logreg.score(X_test,y_test)

print("Testing Data Accuracy by Logistics Regression is : " , logreg_test)
```

Training Data Accuracy by Logistics Regression Algorithm is : 0.8280373831775701
Testing Data Accuracy by Logistics Regression is : 0.8208955223880597

```
In [94]: # Evaluate the model on the test data
score = logreg.score(X_test, y_test)
print("Accuracy of Logistic Regression is : ",score)
```

Accuracy of Logistic Regression is : 0.8208955223880597

```
In [95]: # calculating the mean squared error
mse = np.mean((y_test - y_pred)**2, axis = None)
print("MSE :", mse)

# Calculating the root mean squared error
rmse = np.sqrt(mse)
print("RMSE :", rmse)
```

MSE : 0.1791044776119403
RMSE : 0.42320736951515897

Deciding a Model

```
In [96]: # Usually We Select Only Those Model Which Has Highest Accuracy Among All tho
```

```
In [100]: if tuned_model_test > logreg_test :
           print (" For this Data Highest Accuracy belong to Decision Tree, out of 2
           else:
           print(" For this Highest Accuracy belong Data Logistics Regression, out of
```

For this Highest Accuracy belong Data Logistics Regression, out of 4 model w
ith accuracy of 0.8208955223880597

```
In [81]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```